

A three-layer diagnostic for institutional voice in LLM-generated news
(NYT China coverage; GPT-4o-mini, Gemini-2.5-Flash, Claude-Haiku-4.5; n = 40/desk)

Increasing resistance to prompt engineering



Reproduced

reproduced by default

Entity & topic vocabulary

All 3 models hit desk-specific lexicons spontaneously
(Business fidelity $\approx$ NYT; Foreign fidelity positive).



Recoverable

recovered by prompting

Business-desk affective baseline

Prompt scaffolding (V.E) recovers the Business tone —
Claude reaches NYT (n.s., $p = 0.23$).



Miscalibrated

survives prompting

Foreign-desk affective baseline

Stays more positive than NYT across all 3 families, even under
the strongest prompt. Foreign > Business in 12/12 cells.